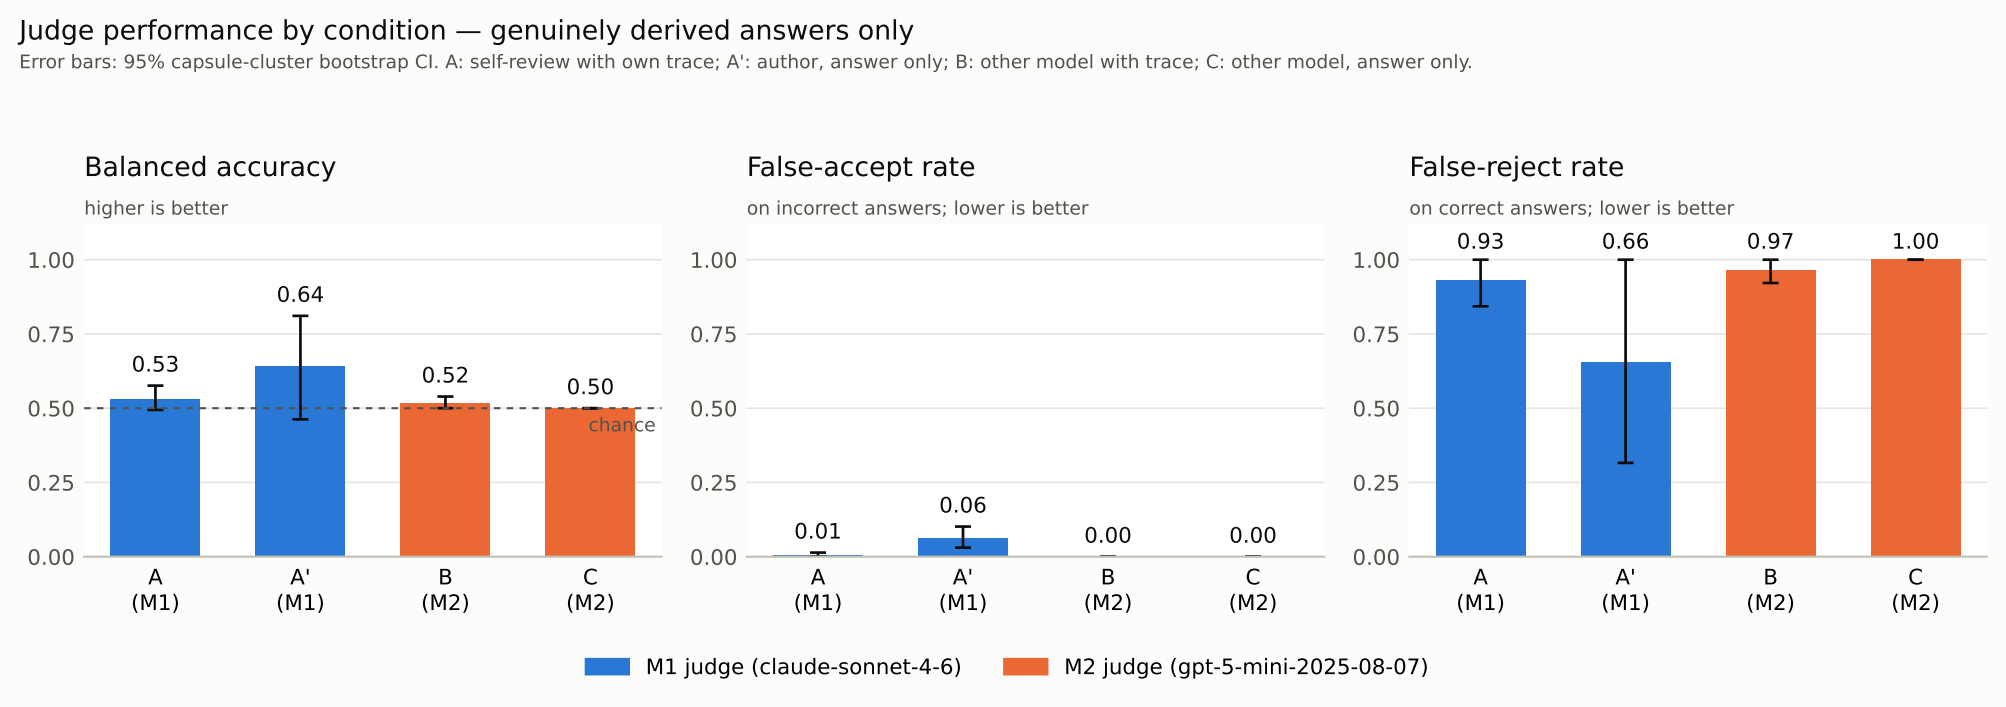
Judge performance by condition — genuinely derived answers only
Error bars: 95% capsule-cluster bootstrap CI. A: self-review with own trace; A': author, answer only; B: other model with trace; C: other model, answer only.
Balanced accuracy
higher is better
1.00
0.75
0.50
0.25
0.00
0.53
0.64
0.52
0.50
chance
A (M1)
A' (M1)
B (M2)
C (M2)
False-accept rate
on incorrect answers; lower is better
0.01
0.06
0.00
0.00
False-reject rate
on correct answers; lower is better
0.93
0.66
0.97
1.00
M1 judge (claude-sonnet-4-6)
M2 judge (gpt-5-mini-2025-08-07)